



Console Games Database

ISM6218.901S22 Advanced Database Management

Group 2

Shailendra Singh

Sherry Ingram

Abhishanth Penta

Kshitiz Kharel

Executive Summary

Introduction

A console game is a type of interactive multimedia software that uses a video game console to provide an interactive multimedia experience via a television or other display device. The game console generally consists of a handheld control device (although some use cameras to monitor user movements) and a computer that runs the game's software. The global console games market is about \$26.8b in 2018 with key vendors in the market like Microsoft, Sony, and Nintendo.

Our group decided to create a database that serves as the backend for any application or services, which can leverage it to provide product recommendations (videogames) based on key attributes such as rating, genre, console compatibility, and so on. Our goal is to set up a working database that can also be used for data-science tasks such as predicting success of a new game based on analyzing similar games' information stored in our database.

Project Topic Areas and Weights

| Topic Area | Description | Points |
|--------------------------------------|--|------------------------------|
| Database Design | This part includes a logical database design (for the relational model), using normalization to control redundancy and integrity constraints for data quality. | Weight: 30 Range: 20 - 30 |
| Query Writing | This part demonstrates practical examples using SQL queries on the database for real-world scenarios or data request within an application/service. | Weight: 25 Range: 20 - 30 |
| Performance Tuning | This section applies some performance tuning concepts such as indexing, etc., that improves the performance of the query executions. | Weight: 20 Range: 20 - 30 |
| Data Visualization and ML Algorithms | This section attempts to demonstrate to use visualizations and data mining algorithms to analyze the data. | Weight: 25 Range: 10 - 40 |

Database Design

Section 1.1: Data Generation and Loading

Data Collection




The data was collected from two sources.

Games Data

The games data primarily comes from *Online Game System Repository* (<https://www.gamesdatabase.org/>). The following screenshot shows one such page:



The screenshot shows a web browser window with the URL `gamesdatabase.org/list.aspx?DM=0&category=racing&sort=Game`. The page displays a table of racing games. The table has columns for Game, System, Publisher, Developer, Category, and Year. The first three rows of data are visible:

| Game | System | Publisher | Developer | Category | Year |
|--|--|---|--|----------|------|
|  0 Gravity Y3030 |  Microsoft Xbox Live Arcade |  Evolving Software | Racing | | |
|  007: Everything or Nothing |  Nintendo Game Boy Advance |  Electronic Arts |  Griptonite Games | Racing | 2003 |
|  007: From Russia with Love |  Sony PSP |  Electronic Arts |  Electronic Arts | Racing | 2006 |

Since the webpage has a defined structure and predictable formatting, we used python to web-scrap this data. We took only those categories that has more than 1000 games.

Users Data

For Users data, we used *Kaggle* dataset (<https://www.kaggle.com/datasets/nathanlauga/nba-games>). This dataset contains games and players information of NBA.

Data Cleaning

We used *Python* to clean the data and convert them into required CSV files. The following script shows the csv creation for games data.

```
import os

class Record:
    def __init__(self, line_txt):
        items = line_txt.split('\t')
        self.game = items[1]
        self.system = items[3]
        self.publisher = items[5]
        self.developer = items[7]
        self.category = items[8]
        self.year = items[9]
```

```

        if not (self.game and self.system and self.publisher and self.developer and
self.category):
            raise ValueError

    def __str__(self):
        return "{}{}, {}, {}, {}, {}, {}".format(self.game, self.system, self.publisher,
self.developer, self.category,
                                                    self.year)

def get_records(file_name):
    records = []
    with open(file_name) as f:
        for line in f.readlines():
            if len(line.split('\t')) == 10:
                try:
                    rec = Record(line)
                    records.append(rec)
                except ValueError:
                    pass
    return records

output_dir = r"D:\Practice\PythonWorkspace\DBMS Games Data Archive"
output_file = os.path.join(r'D:\Practice\PythonWorkspace\Notebooks',
"games_data.csv")
with open(output_file, 'w') as out:
    header = Record("    Game        System    Publisher    Developer    Category
Year\n")
    out.write(str(header))
    for dirName, subdirList, fileList in os.walk(output_dir):
        for fl in fileList:
            res = get_records(os.path.join(dirName, fl))
            res = [str(x) for x in res]
            out.writelines(res)

print("Done")

```

This script read all the raw files and collated them into one final CSV file. The final csv file looks something like shown below:

```

> games_data_cleaned = pd.read_csv(r"D:\Practice\PythonWorkspace\Notebooks\games_data_cleaned.csv")
print("(Row,Col) = {}".format(games_data_cleaned.shape))
games_data_cleaned.head()
[20] ✓ 1.1s Python
... (Row,Col) = (24931, 6)
</>

```

| | Game | System | Publisher | Developer | Category | Year |
|---|----------------------------|--------------------|-----------------|-----------------|----------|------|
| 0 | [PROTOTYPE] | Microsoft Xbox 360 | Activision | Radical | Action | NaN |
| 1 | 007: Agent Under Fire | Microsoft Xbox | Electronic Arts | Electronic Arts | Action | NaN |
| 2 | 007: Agent Under Fire | Nintendo GameCube | Electronic Arts | Electronic Arts | Action | NaN |
| 3 | 007: Everything or Nothing | Nintendo GameCube | Electronic Arts | Electronic Arts | Action | NaN |
| 4 | 007: Everything or Nothing | Microsoft Xbox | Electronic Arts | Electronic Arts | Action | NaN |

The Users data was cleaned in similar style. The following code shows the cleaning process for all the tables in the ER diagram (shown later in the document).

Load Libraries

```
import pandas as pd
import numpy as np
import random
```

Read NBA Kaggle Data Files

```
players = pd.read_csv('games_data/players.csv', on_bad_lines='skip')
games = pd.read_csv('games_data/games.csv', on_bad_lines='skip')
games_details = pd.read_csv('games_data/games_details.csv', on_bad_lines='skip')
ranking = pd.read_csv('games_data/ranking.csv', on_bad_lines='skip')
teams = pd.read_csv('games_data/teams.csv', on_bad_lines='skip')
console = pd.read_csv('games_data/games_data_cleaned.csv')
```

Create User Table

```
# Create Gamer Table
def get_tag(fname, lname=None):
    try:
        if lname:
            return fname[:3]+lname[:3]
    except:
        return fname[:3] + fname[:-3:-1]

    return fname[:3]+fname[:-3:-1]

gamer_table=games_details[['PLAYER_ID', 'PLAYER_NAME', 'NICKNAME', 'TEAM_CITY', 'PF', 'FG_PCT']].drop_duplicates().sort_values(by='PLAYER_ID')
gamer_table=gamer_table.groupby('PLAYER_ID').first()
gamer_table=gamer_table.reset_index()
gamer_table['FirstName']=gamer_table['PLAYER_NAME'].str.split(' ').str[0]
gamer_table['LastName']=gamer_table['PLAYER_NAME'].str[1]
gamer_table['FirstName']=gamer_table['FirstName'].str[0]

gamer_table['NICKNAME'] = gamer_table.apply(lambda x: get_tag(x.FirstName, x.LastName), axis=1)
gamer_table['GamerScore']=round(gamer_table['FG_PCT']*100)
gamer_table['Age'] = np.random.randint(13, 45, gamer_table.shape[0])
gamer_table[gamer_table['GamerScore'].isnull()] =
gamer_table[gamer_table['GamerScore'].isnull()]

random.randint(11,99))
gamer_table['GamerScore']=gamer_table['GamerScore'].astype('int32')
gamer_table.rename(columns={'PLAYER_ID':'UID', 'NICKNAME':'GamerTag', 'TEAM_CITY':'City'}, inplace=True)
gamer_table =
gamer_table[['UID', 'FirstName', 'LastName', 'GamerTag', 'City', 'Age', 'GamerScore']].drop_duplicates().sort_values(by='UID')
```

Use standard data to get state from city names

```

cities = pd.read_csv('games_data/uscities.csv', on_bad_lines='skip')
cities = cities[['city', 'state_name']]
cities['City'] = cities['city']
cities = cities[['City', 'state_name']]
gamer_table=pd.merge(
    gamer_table,
    cities,
    on="City"
)
gamer_table=gamer_table.groupby('UID').first()
gamer_table=gamer_table.reset_index()
gamer_table['GamerTag']=gamer_table['GamerTag'].apply(str.lower)
gamer_table.rename(columns={'state_name':'State'}, inplace=True)

```

Further clean games data to match with project plan ER diagram

```

# Remove useless system values.
console=console[console['System']!='4']

# Fix System Column
console['System2']=console['System'].str.split(' ')
console['System2']=console['System2'].str[0]
console['System']=console['System2']
console=console.drop('System2',axis=1)
console.head()

# Fix Year Column
console['Year']=console['Year'].apply(pd.to_numeric)

# Create Categorical Column Codes
#-----
# Set all desired cols as category
console['CategoryCode'] = console['Category'].astype('category')
console['SystemCode'] = console['System'].astype('category')
console['PublisherCode'] = console['Publisher'].astype('category')
# Select all category columns and apply cat.codes attribute
cat_cols = console.select_dtypes(['category']).columns
console[cat_cols]=console[cat_cols].apply(lambda x : x.cat.codes+1)

```

Create Games Table

```

# Create Games Table
games_temp = games[['GAME_ID', 'REB_home', 'FT_PCT_away']].drop_duplicates()
games_temp.columns = ('GameId', 'Price', 'Rating')
console_temp = console[['Game', 'CategoryCode']].drop_duplicates()
console_temp = console_temp.reset_index()
games_temp = games_temp[:console_temp.shape[0]]
games_temp=games_temp.reset_index()
games_temp['index']=np.arange(1,games_temp.shape[0]+1,1)
console_temp['index']=np.arange(1,games_temp.shape[0]+1,1)
games_table = pd.merge(games_temp, console_temp,on="index")
games_table['Rating']=games_table['Rating']*5
games_table['Rating']=games_table['Rating'].round(2)
games_table = games_table[['GameId', 'Game', 'CategoryCode', 'Price', 'Rating']]
games_table.rename(columns={'Game':'Name', 'CategoryCode':'GenreId'}, inplace=True)
games_table['AgeRating']=np.random.randint(1, 4, games_table.shape[0])

```

Create Mapping table to preserve 3rd Normal Form

```
# Create GameConsoleAvailability
game_console_table = pd.merge(console[['Game', 'SystemCode']],
                              games_table[['Name', 'GameId']],
                              left_on="Game",
                              right_on="Name")
game_console_table = game_console_table[['GameId', 'SystemCode']].drop_duplicates()
game_console_table.columns = ('GameId', 'ConsoleId')
```

Create Dimension Table - Age Ratings

```
# Create Ratings Table
ratings_data = [[1, 'Rated G', 'General audiences – All ages admitted.'],
                [2, 'Rated PG', 'Parental guidance suggested – Some material may not
be suitable for pre-teenagers.'],
                [3, 'Rated R', 'Restricted – Under 17 requires accompanying parent or
adult guardian.'],
                [4, 'Rated X', 'No one under 17 admitted.']]

ratings_table = pd.DataFrame(ratings_data, columns = ['AgeRatingId',
'AgeRatingName', 'AgeRatingDescription'])
```

Create Console Table

```
# Create Console Table
console_table =
console[['System', 'SystemCode']].drop_duplicates().sort_values(by='SystemCode')
console_table = console_table.reset_index(drop=True)
console_table.rename(columns={'SystemCode': 'ConsoleId', 'System': 'ConsoleName'},
inplace=True)
console_table = console_table[['ConsoleId', 'ConsoleName']]
```

Create Genre Table

```
# Create Genre Table
genre_table =
console[['Category', 'CategoryCode']].drop_duplicates().sort_values(by='CategoryCode'
)
genre_table = genre_table.reset_index(drop=True)
genre_table.rename(columns={'Category': 'Genre', 'CategoryCode': 'GenreId'},
inplace=True)
genre_table = genre_table[['GenreId', 'Genre']]
```

Create Subscription Table

```
# Create Subscription Table
games_details_temp = games_details[['GAME_ID', 'PLAYER_ID']]
subs_table =
pd.merge(games_details_temp, gamer_table[['UID', 'GamerTag']], right_on="UID", left_on="
PLAYER_ID")
subs_table = subs_table[['UID', 'GamerTag', 'GAME_ID']]
subs_table =
pd.merge(subs_table, games_table[['GameId']], right_on="GameId", left_on="GAME_ID")
```



```
subs_table = subs_table[['UID', 'GamerTag', 'GameId']]
subs_table['SubscriptionId'] = np.random.randint(1, 3, subs_table.shape[0])
```

Create Mapping table to preserve 3rd Normal Form

```
# Create Game Subscription Type Table
subscription_data = [[1, 'Free'], [2, 'Freemium'], [3, 'Paid']]
game_sub_type_table = pd.DataFrame(subscription_data, columns = ['SubscriptionId',
'SubscriptionName'])
```

Export respective tables to their CSV counterparts

```
from os import path
output_path =
r"C:\Users\Shail\Documents\PythonWorkspace\Notebooks\games_data\output"

# 1. Users Table
gamer_table.to_csv(path.join(output_path, "users.csv"), sep=',', encoding='utf-
8', index=False)
# 2. Games Table
games_table.to_csv(path.join(output_path, "games.csv"), sep=',', encoding='utf-
8', index=False)
# 3. Genre Table
genre_table.to_csv(path.join(output_path, "genre.csv"), sep=',', encoding='utf-
8', index=False)
# 4. Game Console Table
game_console_table.to_csv(path.join(output_path, "game_console.csv"), sep=',',
encoding='utf-8', index=False)
# 5. Age Ratings Table
ratings_table.to_csv(path.join(output_path, "age_ratings.csv"), sep=',',
encoding='utf-8', index=False)
# 6. Console Table
console_table.to_csv(path.join(output_path, "consoles.csv"), sep=',', encoding='utf-
8', index=False)
# 7. Game Subscription Type Table
game_sub_type_table.to_csv(path.join(output_path, "subscription_types.csv"), sep=',',
encoding='utf-8', index=False)
# 8. Subscription Table
subs_table.to_csv(path.join(output_path, "subscriptions.csv"), sep=',',
encoding='utf-8', index=False)
```

The last segment created following tables. The first line shows Row and Column count and the second shows table itself.

Age Rating

... (Row,Col) = (4, 3)

| | AgeRatingId | AgeRatingName | AgeRatingDescription |
|---|-------------|---------------|---|
| 0 | 1 | Rated G | General audiences – All ages admitted. |
| 1 | 2 | Rated PG | Parental guidance suggested – Some material ma... |
| 2 | 3 | Rated R | Restricted – Under 17 requires accompanying pa... |
| 3 | 4 | Rated X | No one under 17 admitted. |

Consoles

... (Row,Col) = (35, 2)

| | ConsoleId | ConsoleName |
|---|-----------|-------------|
| 0 | 1 | Acorn |
| 1 | 2 | Amstrad |
| 2 | 3 | Apple |
| 3 | 4 | Arcade |
| 4 | 5 | Atari |

Game Consoles

... (Row,Col) = (23354, 2)

| | Gameld | ConsoleId |
|---|----------|-----------|
| 0 | 22101005 | 20 |
| 1 | 22101006 | 20 |
| 2 | 22101006 | 22 |
| 3 | 22101007 | 22 |
| 4 | 41000172 | 22 |

Games

... (Row,Col) = (24931, 6)

| | Game | System | Publisher | Developer | Category | Year |
|---|----------------------------|--------------------|-----------------|-----------------|----------|------|
| 0 | [PROTOTYPE] | Microsoft Xbox 360 | Activision | Radical | Action | NaN |
| 1 | 007: Agent Under Fire | Microsoft Xbox | Electronic Arts | Electronic Arts | Action | NaN |
| 2 | 007: Agent Under Fire | Nintendo GameCube | Electronic Arts | Electronic Arts | Action | NaN |
| 3 | 007: Everything or Nothing | Nintendo GameCube | Electronic Arts | Electronic Arts | Action | NaN |
| 4 | 007: Everything or Nothing | Microsoft Xbox | Electronic Arts | Electronic Arts | Action | NaN |

Genre

... (Row,Col) = (20, 2)

| | GenreId | Genre |
|---|---------|---------------|
| 0 | 1 | Action |
| 1 | 2 | Adventure |
| 2 | 3 | Application |
| 3 | 4 | Arcade |
| 4 | 5 | Ball & Paddle |

Subscription Types

... (Row,Col) = (3, 2)

| | SubscriptionId | SubscriptionName |
|---|----------------|------------------|
| 0 | 1 | Free |
| 1 | 2 | Freemium |
| 2 | 3 | Paid |

Subscriptions

... (Row,Col) = (317950, 4)

| | UID | GamerTag | GameId | SubscriptionId |
|---|---------|----------|----------|----------------|
| 0 | 1627736 | malbea | 22101005 | 2 |
| 1 | 1626156 | d'arus | 22101005 | 2 |
| 2 | 1627752 | taupri | 22101005 | 1 |
| 3 | 1630233 | natkni | 22101005 | 1 |
| 4 | 1629130 | dunrob | 22101005 | 2 |

Users

... (Row,Col) = (2151, 8)

</>

| | UID | FirstName | LastName | GamerTag | City | Age | GamerScore | State |
|---|-----|-----------|------------|----------|--------------|-----|------------|----------------------|
| 0 | 15 | Eric | Piatkowski | eripia | Chicago | 44 | 67 | Illinois |
| 1 | 42 | Monty | Williams | monwil | Philadelphia | 39 | 25 | Pennsylvania |
| 2 | 43 | Chris | Whitney | chrwhi | Washington | 19 | 60 | District of Columbia |
| 3 | 56 | Gary | Payton | garpay | Miami | 24 | 40 | Florida |
| 4 | 57 | Doug | Christie | douchr | Dallas | 23 | 67 | Texas |

Data Loading

The respective CSV files were loaded using Oracle SQL Developers Import Wizard.

We first created all the tables, following SQL Shows CREATE Statements:

```
CREATE TABLE DB1SMOKE.AGE_RATINGS
(
  AgeRatingId INT,
  AgeRatingName VARCHAR2(10),
  AgeRatingDescription VARCHAR2(100)
);
```

```
CREATE TABLE DB1SMOKE.CONSOLES
(
  ConsoleId INT,
  ConsoleName VARCHAR2(20)
);
```

```
CREATE TABLE DB1SMOKE.GAME_CONSOLES
(
  GameId INT,
  ConsoleId INT
);
```

```
CREATE TABLE DB1SMOKE.GAMES
(
  GameId INT,
  GameName VARCHAR2(100),
  GenreId INT,
  Price FLOAT,
  Rating FLOAT,
  AgeRating FLOAT
);
```

```
CREATE TABLE DB1SMOKE.GENRE
(
  GenreId INT,
  Genre VARCHAR2(20)
);
```

```
CREATE TABLE DB1SMOKE.SUBSCRIPTION_TYPES
```

```
(
SubscriptionId INT,
SubscriptionName VARCHAR2(10)
);

CREATE TABLE DB1SMOKE.SUBSCRIPTIONS
(
UserId INT,
GamerTag VARCHAR2(10),
GameId INT,
SubscriptionId INT
);

CREATE TABLE DB1SMOKE.USERS
(
UserId INT,
FirstName VARCHAR2(20),
LastName VARCHAR2(20),
GamerTag VARCHAR2(10),
City VARCHAR2(20),
Age INT,
GamerScore INT,
State VARCHAR2(30)
);
```

Section 1.2 Data Integrity

The integrity constraints necessary to help ensure data quality should be included in the design section. We implemented following constraints to ensure data integrity:

Primary Key Constraints

The following table are dimension tables in our database design. So, we need to create a Primary Keys on columns that are to be used as a Foreign Key in other tables.

```
ALTER TABLE AGE_RATINGS ADD CONSTRAINT PK_AGE_RATING_ID PRIMARY KEY
(AGERATINGID);

ALTER TABLE GENRE ADD CONSTRAINT PK_GENRE_ID PRIMARY KEY (GENREID);

ALTER TABLE SUBSCRIPTION_TYPES ADD CONSTRAINT PK_SUBSCRIPTION_ID PRIMARY KEY
(SUBSCRIPTIONID);

ALTER TABLE CONSOLES ADD CONSTRAINT PK_CONSOLE_ID PRIMARY KEY (CONSOLEID);
```

The following columns are unique identifiers in their respectable tables. This ensures that there are no duplicate records.

```
ALTER TABLE GAMES ADD CONSTRAINT PK_GAME_ID PRIMARY KEY (GAMEID);
```

```
ALTER TABLE USERS ADD CONSTRAINT PK_USER_ID PRIMARY KEY (USERID);
```

Foreign Key Constraints

We added foreign keys in our tables to ensure referential integrity across our database design. This constraint helps us to define the relationship between the tables.

For example, a foreign key reference in GAMES table to the GENRE table shows that the genre information can be acquired from the target table. This also helps with saving storage space as large string values are avoided being repeated and are replaced by their much shorter id counterparts.

```
ADD FOREIGN KEY constraints.
```

```
ALTER TABLE GAMES
```

```
ADD CONSTRAINT FK_GENRE_ID
```

```
FOREIGN KEY (GENREID) REFERENCES GENRE(GENREID);
```

```
ALTER TABLE SUBSCRIPTIONS
```

```
ADD CONSTRAINT FK_SUBSCRIPTION_ID
```

```
FOREIGN KEY (SUBSCRIPTIONID) REFERENCES SUBSCRIPTION_TYPES(SUBSCRIPTIONID);
```

```
ALTER TABLE GAME_CONSOLES
```

```
ADD CONSTRAINT FK_CONSOLE_GAME_ID
```

```
FOREIGN KEY (GAMEID) REFERENCES GAMES(GAMEID);
```

```
ALTER TABLE GAME_CONSOLES
```

```
ADD CONSTRAINT FK_CONSOLE_ID
```

```
FOREIGN KEY (CONSOLEID) REFERENCES CONSOLES(CONSOLEID);
```

```
ALTER TABLE GAMES
```

```
ADD CONSTRAINT FK_AGE_RATING
```

```
FOREIGN KEY (AGERATING) REFERENCES AGE_RATINGS(AGERATINGID);
```

The foreign key also has one more advantage. It ensures the primary key constraint in the target table, and when used with index, it can significantly improve the performance of any operation that's pulling data from very large tables linked by indexed foreign keys. The subscription table has more than 300K records, and it would be useful to have foreign keys to ensure referential integrity as well as data integrity.

```
ALTER TABLE SUBSCRIPTIONS
```

```
ADD CONSTRAINT FK_GAME_ID
```

```
FOREIGN KEY (GAMEID) REFERENCES GAMES(GAMEID);
```

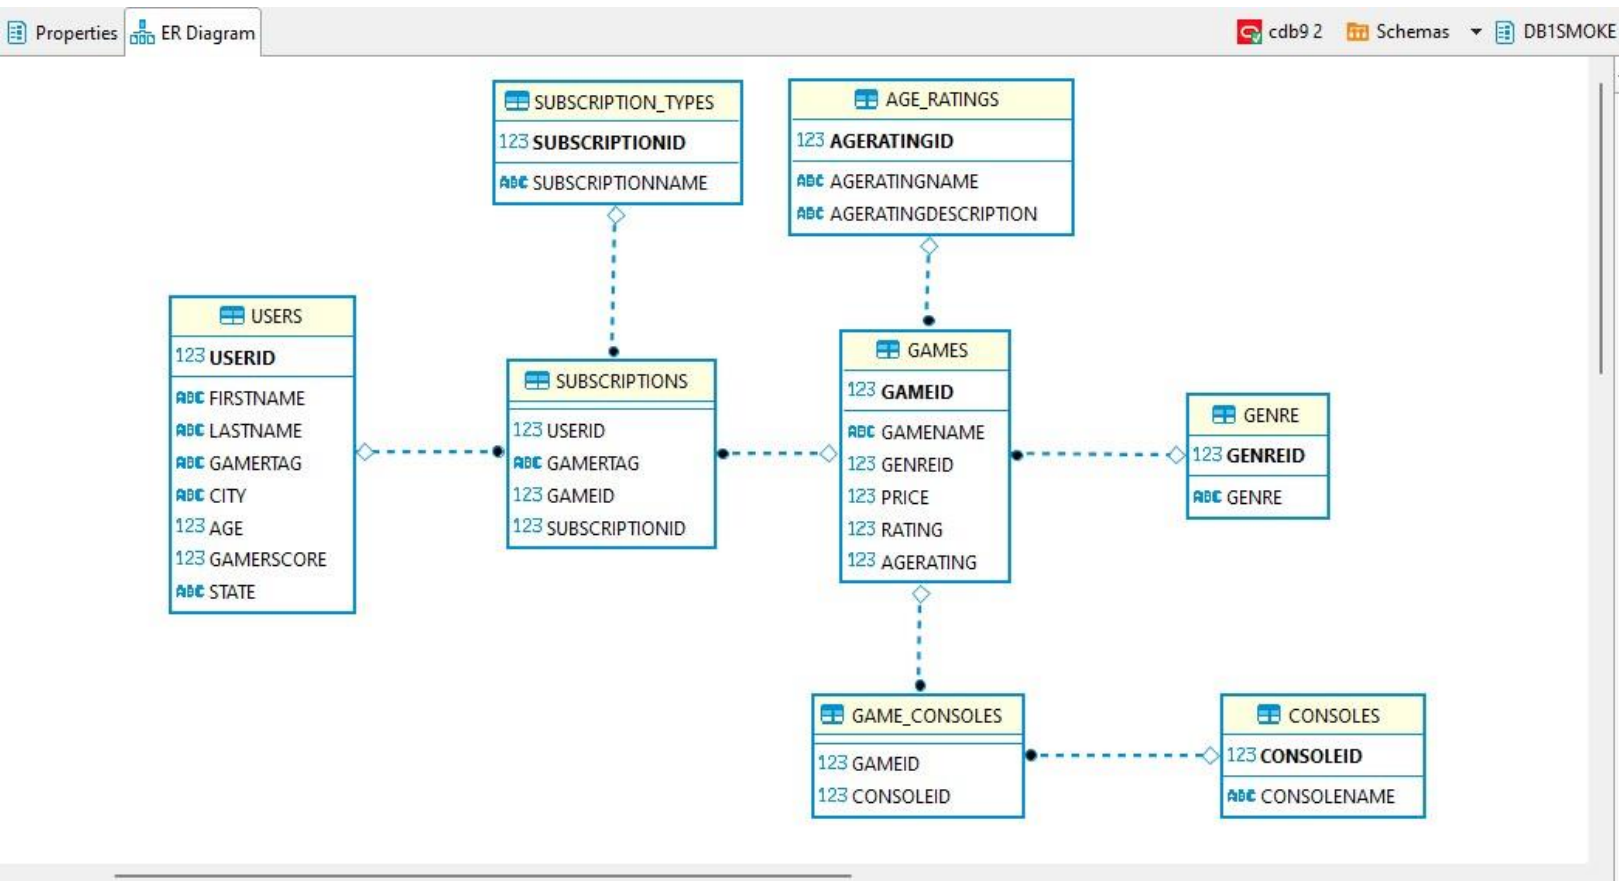
```
ALTER TABLE SUBSCRIPTIONS
```

```
ADD CONSTRAINT FK_USER_ID
```

```
FOREIGN KEY (USERID) REFERENCES USERS(USERID);
```

Section 1.3 E-R Diagram

Our final design looks as shown in below E-R Diagram. Due to the foreign key constraints, we can identify relationships among the tables. The primary keys are shown in bold and the foreign keys can be identified by the diamond symbol in the connecting lines.



Query Writing

This section demonstrates the practical examples where and how the data can be used. We provide several query examples that highlight the types of questions that can be answered using our database.

We present this in question-query format to demonstrate the practical usage.

Q-01. Display Minimum, Maximum, and Average Price per Game Genre. This can be used, for example, to analyze financial aspects by genre and investing appropriately.

```
SELECT
    g2.GENRE,
    MIN(g.PRICE) AS MIN_PRICE,
    MAX(g.PRICE) AS MAX_PRICE,
    AVG(g.PRICE) AS AVG_PRICE
FROM
    DB1SMOKE.GAMES g
INNER JOIN DB1SMOKE.GENRE g2
ON
    G.GENREID = G2.GENREID
GROUP BY
    g2.GENRE
```

| | ABC GENRE | 123 MIN_PRICE | 123 MAX_PRICE | 123 AVG_PRICE |
|---|---------------|---------------|---------------|---------------|
| 1 | Family | 26 | 62 | 42.7108433735 |
| 2 | Racing | 21 | 66 | 41.7217465753 |
| 3 | Ball & Paddle | 35 | 56 | 45.75 |
| 4 | Breakout | 30 | 55 | 41.7358490566 |
| 5 | Educational | 24 | 59 | 43.54375 |

Q-02. Find which cities would be good for hosting gaming competitions.

1. Display the city's name as well as how many games are played by the users are in the city.
2. Show cities with the most counts first.

```
SELECT
    u.CITY,
    COUNT(s.GAMEID) AS GAME_COUNTS
FROM
    DB1SMOKE.USERS u
INNER JOIN DB1SMOKE.SUBSCRIPTIONS s
ON
    u.USERID = s.USERID
GROUP BY
    u.CITY
```


ORDER BY COUNT(s.GAMEID) DESC

| | ABC CITY | 123 GAME_COUNTS |
|---|-------------|-----------------|
| 1 | Los Angeles | 22,148 |
| 2 | Dallas | 17,650 |
| 3 | San Antonio | 17,080 |
| 4 | Indiana | 16,113 |
| 5 | Houston | 14,616 |

Q-03. Display all Games that contains “Need for Speed” somewhere in the name. Show the Games in decreasing order of their price.

```
SELECT
    g.GAMEID ,
    g.GAMENAME ,
    g.PRICE
FROM
    DB1SMOKE.GAMES g
WHERE
    g.GAMENAME LIKE '%Need for Speed%'
ORDER BY
    g.PRICE DESC
```

| | 123 GAMEID | ABC GAMENAME | 123 PRICE |
|---|------------|---------------------------------|-----------|
| 1 | 22,000,049 | Need for Speed Hot Pursuit | 54 |
| 2 | 21,000,576 | Need for Speed: Most Wanted | 53 |
| 3 | 21,000,581 | Need for Speed: V-Rally 2 | 52 |
| 4 | 21,000,596 | Need for Speed III: Hot Pursuit | 50 |
| 5 | 21,000,580 | Need for Speed: Shift | 49 |

Q-04. Label the cost category of a game based on its price. For each game display the game's name, price, and a textual label describing the cost category of the game.

The label should be

- "Very High" for a price more than 60
- "High" for a price range of 45 to 60
- "Average" for a price of 30 to 45, and
- "Low" for a price to be less than 30.

```
SELECT
    g.GAMENAME,
    g.PRICE,
    CASE
        WHEN g.PRICE >= 60 THEN 'Very High'
        WHEN g.PRICE >= 45 AND g.PRICE < 60 THEN 'High'
        WHEN g.PRICE >= 30 AND g.PRICE < 45 THEN 'Average'
        ELSE 'Low'
```

```

END AS "Cost Category"
FROM
    DB1SMOKE.GAMES g;

```

| | ABC GAMENAME | 123 PRICE | ABC Cost Category |
|----|------------------------------|-----------|-------------------|
| 1 | Army Men: Air Attack 2 | 50 | High |
| 2 | Army Men: Sarge's War | 52 | High |
| 3 | As Aventuras da TV Colosso | 45 | High |
| 4 | Asmik-kun Land | 38 | Average |
| 5 | Assassin | 55 | High |
| 6 | Assassin Special Edition | 49 | High |
| 7 | Assassin's Creed | 54 | High |
| 8 | Assassin's Creed Brotherhood | 43 | Average |
| 9 | Assassin's Creed II | 50 | High |
| 10 | Assassin's Creed Revelations | 40 | Average |

Q-05. Find the top 10 most popular consoles. A console is popular when the number of games available exceed 100. Show the results with the most user count per console first.

```

WITH GameConsole AS
(
    SELECT
        c.CONSOLEID,
        c.CONSOLENAME ,
        COUNT(gc.GAMEID) AS GAME_COUNT
    FROM
        DB1SMOKE.CONSOLES c
    INNER JOIN DB1SMOKE.GAME_CONSOLES gc
    ON
        c.CONSOLEID = gc.CONSOLEID
    GROUP BY
        c.CONSOLEID,
        c.CONSOLENAME
    HAVING
        COUNT(gc.GAMEID) > 100
)
SELECT
    c.CONSOLENAME,
    COUNT(u.USERID) AS "USAGE COUNT"
FROM
    GameConsole c
INNER JOIN DB1SMOKE.GAME_CONSOLES gc
ON
    c.CONSOLEID = gc.CONSOLEID
INNER JOIN DB1SMOKE.SUBSCRIPTIONS s
ON
    s.GAMEID = gc.GAMEID

```

```

INNER JOIN DB1SMOKE.USERS u
ON
    u.USERID = s.USERID
GROUP BY
    c.CONSOLENAME
ORDER BY
    COUNT(u.USERID) DESC;

```

| | ABC CONSOLENAME | 123 USAGE COUNT |
|----|-----------------|-----------------|
| 1 | Arcade | 69,019 |
| 2 | Commodore | 65,939 |
| 3 | Nintendo | 64,287 |
| 4 | Atari | 46,297 |
| 5 | Microsoft | 44,660 |
| 6 | Sega | 41,162 |
| 7 | Sony | 36,830 |
| 8 | Sinclair | 30,699 |
| 9 | Amstrad | 22,841 |
| 10 | Valve | 16,239 |

Q-06. Find the top 3 most experienced gamers. A gamer is experienced based on the number of games they have played. Also show which city do they belong to.

```

SELECT
    u.FIRSTNAME ,
    u.LASTNAME ,
    u.CITY,
    count(s.GAMEID) AS "Games Played"
FROM
    DB1SMOKE.USERS u
INNER JOIN DB1SMOKE.SUBSCRIPTIONS s
ON
    u.USERID = s.USERID
GROUP BY
    u.FIRSTNAME ,
    u.LASTNAME ,
    u.CITY
ORDER BY
    count(s.GAMEID) DESC
FETCH NEXT 3 ROWS ONLY

```

| | ABC FIRSTNAME | ABC LASTNAME | ABC CITY | 123 Games Played |
|---|---------------|--------------|-------------|------------------|
| 1 | LeBron | James | Miami | 1,045 |
| 2 | Dwight | Howard | Los Angeles | 1,003 |
| 3 | Carmelo | Anthony | New York | 945 |

Q-07. Find which age rating category has the most played games. Show user count by their age rating category in descending order.

```
SELECT
    ar.AGERATINGNAME,
    COUNT(s.USERID) AS "Usage Count"
FROM
    DB1SMOKE.GAMES g
INNER JOIN DB1SMOKE.SUBSCRIPTIONS s
ON
    g.GAMEID = s.GAMEID
INNER JOIN DB1SMOKE.AGE_RATINGS ar
ON
    g.AGERATING = ar.AGERATINGID
GROUP BY
    ar.AGERATINGNAME
ORDER BY COUNT(s.USERID) DESC ;
```

| | ABC AGERATINGNAME | 123 Usage Count |
|---|-------------------|-----------------|
| 1 | Rated R | 107,434 |
| 2 | Rated G | 105,777 |
| 3 | Rated PG | 104,739 |

Q-08. Find which subscription model is the most popular. The popularity of a subscription model can be judged by the number of users play the games under that subscription type.

```
SELECT
    st.SUBSCRIPTIONNAME,
    COUNT(s.USERID) AS "Usage"
FROM
    DB1SMOKE.SUBSCRIPTIONS s
INNER JOIN DB1SMOKE.SUBSCRIPTION_TYPES st
ON
    s.SUBSCRIPTIONID = st.SUBSCRIPTIONID
GROUP BY
    st.SUBSCRIPTIONNAME
ORDER BY
    COUNT(s.USERID) DESC;
```

| | ABC SUBSCRIPTIONNAME | 123 Usage |
|---|----------------------|-----------|
| 1 | Free | 159,135 |

Q-09. Find all duplicate users. A user is considered duplicate if they have the same first name, last name, city, and state.

```

SELECT
    u.FIRSTNAME,
    u.LASTNAME,
    u.CITY,
    u.STATE,
    COUNT(*) AS FREQ_CNT
FROM
    DB1SMOKE.USERS u
GROUP BY
    u.FIRSTNAME,
    u.LASTNAME,
    u.CITY,
    u.STATE
HAVING
    COUNT(*) > 1;

```

| | ABC FIRSTNAME | ABC LASTNAME | ABC CITY | ABC STATE | 123 FREQ_CNT |
|---|---------------|--------------|-----------|-----------|--------------|
| 1 | Josh | Akognon | Dallas | Texas | 2 |
| 2 | Tristan | Thompson | Cleveland | Ohio | 2 |
| | | | | | |

Q-10. Find the top 3 most popular games.

```

SELECT
    g.GAMENAME,
    COUNT(s.USERID) AS "Usage"
FROM
    DB1SMOKE.GAMES g
INNER JOIN DB1SMOKE.SUBSCRIPTIONS s
ON
    g.GAMEID = s.GAMEID
GROUP BY
    g.GAMENAME
ORDER BY
    COUNT(s.USERID) DESC
FETCH NEXT 3 ROWS ONLY;

```

| | ABC GAMENAME | 123 Usage |
|---|-----------------|-----------|
| 1 | Missile Command | 87 |
| 2 | Warlords | 87 |
| 3 | Star Wars | 80 |

Performance Tuning

Overview

Since this database contains elements of user and product information, it can grow quite huge overtime when new data is introduced. Database performance tuning will allow us to maximize resource utilization so that any application or service that runs on top of it will benefit from the critical database operations. This is important because even relatively minor database-related performance issues can impact the entire operation.

Indexing Strategy

Indexing is a very important strategy that can improve complex query performance involving joins, aggregation, etc. The benefit of indexes lies in the fact that it provides faster access to data for operations that return small number of a table's rows.

In our database, we have identified several columns that we suspect will be used more often than others. It is also a good idea to identify columns that are directly used in join conditions. Following is the description of all the indexes we created:

| Column Name | Source Table | Reason |
|----------------|---------------|---|
| GameId | Game_Consoles | Game Console table is a mapping table. So, this will be the most frequently used table for connecting game id to console id. Console Id is also a frequently used column which has limited values since it is referencing a dimension table. |
| ConsoleId | | |
| GenreId | Games | These two columns are the primary keys of their respective dimension tables. Hence, they will offer value in filtering queries. |
| AgeRating | | |
| UserId | Subscriptions | Subscription table is a central table that has many frequently used columns in join conditions. |
| GameId | | |
| SubscriptionId | | |
| GamerTag | | |
| GamerTag | Users | GamerTag works as a username so instead of finding a user by his/her name, this can act as pseudo-unique identifier. |

Following lines of code shows the SQL statements for creating indexes on above mentioned table-column pairs:

```
-- Creating Indexes for Performance Tuining
```

```
CREATE INDEX IDX_GAMECONSOLE_GAMEID  
ON  
GAME_CONSOLES(GAMEID);
```

```
CREATE INDEX IDX_GAMECONSOLE_CONSOLEID  
ON  
GAME_CONSOLES(CONSOLEID);
```

```
CREATE INDEX IDX_GAMES_GENREID  
ON  
GAMES(GENREID);
```

```
CREATE INDEX IDX_GAMES_AGERATING  
ON  
GAMES(AGERATING);
```

```
CREATE INDEX IDX_SUBSCRIPTIONS_USERID  
ON  
SUBSCRIPTIONS(USERID);
```

```
CREATE INDEX IDX_SUBSCRIPTIONS_GAMEID  
ON  
SUBSCRIPTIONS(GAMEID);
```

```
CREATE INDEX IDX_SUBSCRIPTIONS_SUBSID  
ON  
SUBSCRIPTIONS(SUBSCRIPTIONID);
```

```
CREATE INDEX IDX_SUBSCRIPTIONS_GAMERTAG  
ON  
SUBSCRIPTIONS(GAMERTAG);
```

```
CREATE INDEX IDX_USERS_GAMERTAG  
ON  
USERS(GAMERTAG);
```

DB experiment for index performances

In previous section, we have created indexes on our main or production tables. Before loading the data into Oracle, we also had created the staging tables for backup. In this experiment, we will run a query on staging tables, that have no indexes on them, and on main tables, which have indexes created. We plan to compare the execution plans to see if they help in improving the query performance.

Staging Table Query


```

WITH GameConsole AS
(
SELECT
    c.CONSOLEID,
    c.CONSOLENAME ,
    COUNT(gc.GAMEID) AS GAME_COUNT
FROM
    DB1SMOKE.STG_CONSOLES c
INNER JOIN DB1SMOKE.STG_GAME_CONSOLES gc
ON
    c.CONSOLEID = gc.CONSOLEID
GROUP BY
    c.CONSOLEID,
    c.CONSOLENAME
HAVING
    COUNT(gc.GAMEID) > 100
)
SELECT
    c.CONSOLENAME,
    COUNT(u.USERID) AS "USAGE COUNT"
FROM
    GameConsole c
INNER JOIN DB1SMOKE.STG_GAME_CONSOLES gc
ON
    c.CONSOLEID = gc.CONSOLEID
INNER JOIN DB1SMOKE.STG_SUBSCRIPTIONS s
ON
    s.GAMEID = gc.GAMEID
INNER JOIN DB1SMOKE.STG_USERS u
ON
    u.USERID = s.USERID
GROUP BY
    c.CONSOLENAME
ORDER BY
    COUNT(u.USERID) DESC;

```

Execution Plan Results

| OPERATION | OBJECT_NAME | OPTIONS | CARDINALITY | COST |
|--------------------------|-------------------|----------|-------------|------|
| SELECT STATEMENT | | | | |
| SORT | | | | |
| HASH | | | | |
| HASH JOIN | | ORDER BY | 35 | 418 |
| Access Predicates | | GROUP BY | 35 | 418 |
| U.USERID=S.USERID | | | 627254 | 388 |
| TABLE ACCESS | STG_USERS | FULL | 2151 | 7 |
| HASH JOIN | | | 627254 | 379 |
| Access Predicates | | | | |
| S.GAMEID=GC.GAMEID | | | | |
| HASH JOIN | | | 29359 | 35 |
| Access Predicates | | | | |
| C.CONSOLEID=GC.CONSOLEID | | | | |
| VIEW | | | 44 | 20 |
| FILTER | | | | |
| Filter Predicates | | | | |
| COUNT(gc.GAMEID)>100 | | | | |
| HASH | | GROUP BY | 44 | 20 |
| HASH JOIN | | | 23354 | 18 |
| Access Predicates | | | | |
| C.CONSOLEID=GC.CONSOLEID | | | | |
| TABLE ACCESS | STG_CONSOLES | FULL | 35 | 3 |
| TABLE ACCESS | STG_GAME_CONSOLES | FULL | 23354 | 15 |
| TABLE ACCESS | STG_GAME_CONSOLES | FULL | 23354 | 15 |
| TABLE ACCESS | STG_SUBSCRIPTIONS | FULL | 317950 | 343 |

Now we compare this plan with the one ran on main tables with indexes.

```

WITH GameConsole AS
(
SELECT
    c.CONSOLEID,
    c.CONSOLENAME ,
    COUNT(gc.GAMEID) AS GAME_COUNT
FROM
    DB1SMOKE.CONSOLES c
INNER JOIN DB1SMOKE.GAME_CONSOLES gc
ON
    c.CONSOLEID = gc.CONSOLEID
GROUP BY
    c.CONSOLEID,
    c.CONSOLENAME
HAVING
    COUNT(gc.GAMEID) > 100
)
SELECT
    c.CONSOLENAME,
    COUNT(u.USERID) AS "USAGE COUNT"
FROM
    GameConsole c
INNER JOIN DB1SMOKE.GAME_CONSOLES gc
ON
    c.CONSOLEID = gc.CONSOLEID
INNER JOIN DB1SMOKE.SUBSCRIPTIONS s
ON
    s.GAMEID = gc.GAMEID
INNER JOIN DB1SMOKE.USERS u
ON
    u.USERID = s.USERID

```

GROUP BY
c.CONSOLENAME
ORDER BY
COUNT(u.USERID) DESC;

| OPERATION | OBJECT_NAME | OPTIONS | CARDINALITY | COST |
|---|--------------------------|----------------|-------------|-------|
| SELECT STATEMENT | | | | 2 359 |
| SORT | | ORDER BY | 2 | 359 |
| HASH | | GROUP BY | 2 | 359 |
| HASH JOIN | | | 28512 | 356 |
| Access Predicates S.GAMEID=GC.GAMEID | | | | |
| NESTED LOOPS | | | 28512 | 356 |
| NESTED LOOPS | | | | |
| STATISTICS COLLECTOR | | | | |
| HASH JOIN | | | 1335 | 37 |
| Access Predicates C.CONSOLEID=GC.CONSOLEID | | | | |
| VIEW | | | 2 | 21 |
| FILTER | | | | |
| Filter Predicates COUNT(GC.GAMEID) > 100 | | | | |
| HASH | | GROUP BY | 2 | 21 |
| HASH JC | | | 23354 | 19 |
| Access Predicates C.CONSOLEID=GC.CONSOLEID | | | | |
| TABLE ACCESS | CONSOLES | FULL | 35 | 3 |
| TABLE ACCESS | GAME_CONSOLES | FULL | 23354 | 16 |
| TABLE ACCESS | GAME_CONSOLES | FULL | 23354 | 16 |
| INDEX | IDX_SUBSCRIPTIONS_GAMEID | RANGE SCAN | | |
| Access Predicates S.GAMEID=GC.GAMEID | | | | |
| TABLE ACCESS | SUBSCRIPTIONS | BY INDEX ROWID | 21 | 318 |
| Filter Predicates S.USERID IS NOT NULL | | | | |
| TABLE ACCESS | SUBSCRIPTIONS | FULL | 317950 | 318 |
| Filter Predicates S.USERID IS NOT NULL | | | | |

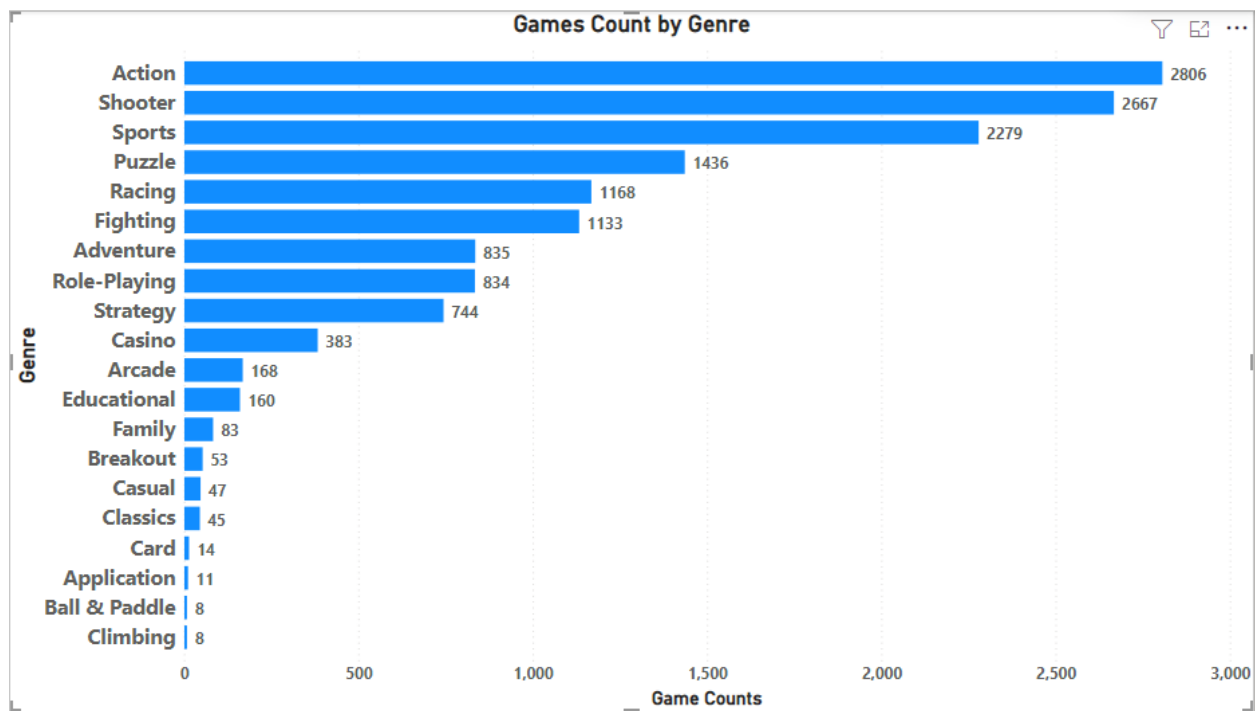
We can observe quite significant difference between the execution plans of the two queries. The last row where the table access occurs has a cost of 318 in indexed table as compared to the higher value of 343 in the non-indexed table. The performance gain occurred due to the BY INDEX ROWID section as pointed by the query coordinator.

Other Topics

Section 1.1: Data Visualization

Data visualization helps to tell stories by curating data into a form easier to understand, highlighting the trends and outliers. A good visualization tells a story, removing the noise from data and highlighting the useful information. We used Microsoft Power BI to visualize the data in our database.

Viz. 01 – Game Counts by Genre



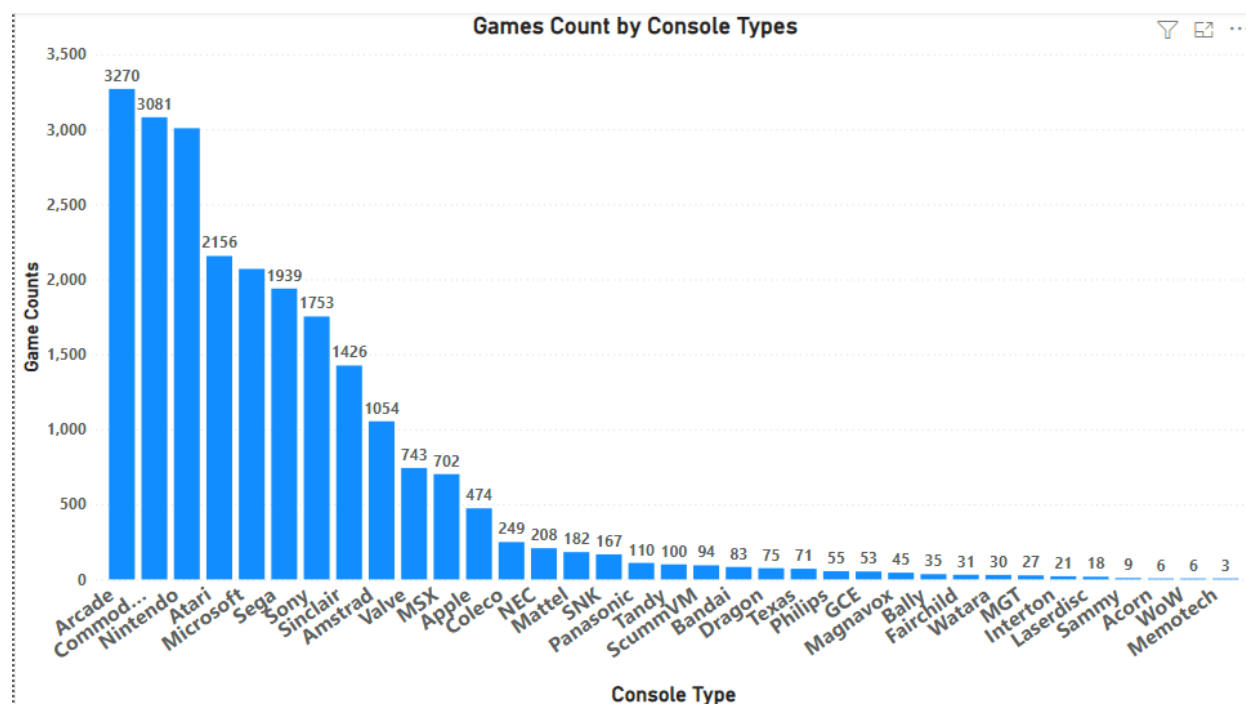
This visualization gives an overall distribution of games by their genres. This also tells what are the types of genres that are most popular among users and hence game developers target games in those categories.

Viz. 02 – Game Counts by Console Types

Similarly, we can observe the most popular consoles based on this data. Please note that this data is taken from a games archive that may not represent the most current trends. Here we see Sony,

Microsoft Xbox, and Nintendo lagging some of the most older console systems. This again can change as we acquire more data.

The best thing about Power BI visuals is that they get updated if the data in their underlying tables change. Following visual helps to view this distribution:



Viz. 03 – Gamer’s Location (Map Visualization)

To visualize location data, Power BI has a map visual, but it requires latitude and longitude data. To achieve that we use a standard dataset which has USA cities geo location data. The CSV looks something like below:

```
cities = pd.read_csv(r'C:\Users\Shail\Downloads\simplemaps_uscities_basicv1.74\uscities.csv', on_bad_lines='skip')
print("(Row,Col) = {}".format(cities.shape))
cities.head()
```

[3] ✓ 0.4s Python

... (Row,Col) = (28338, 17)

| | city | city_ascii | state_id | state_name | county_fips | county_name | lat | lng | population | density | source | military | incorporated |
|---|-------------|-------------|----------|------------|-------------|-------------|---------|-----------|------------|---------|---------|----------|--------------|
| 0 | New York | New York | NY | New York | 36061 | New York | 40.6943 | -73.9249 | 18713220 | 10715 | polygon | False | True |
| 1 | Los Angeles | Los Angeles | CA | California | 6037 | Los Angeles | 34.1139 | -118.4068 | 12750807 | 3276 | polygon | False | True |

We imported this dataset into a table: `DB1SMOKE.CITIES`

On top of it we created following view:

```
CREATE VIEW DB1SMOKE.VW_USER_LOCATION_DATA AS
SELECT
    u.UserId,
    c.county_name AS Region,
    u.City,
    u.State,
    c.state_id AS StateCode,
    u.Age,
    u.GamerScore,
    c.lat AS Latitude,
    c.lng AS Longitude,
    'United States' AS Country,
    'USA' AS CountryCode
FROM
    DB1SMOKE.USERS u
INNER JOIN DB1SMOKE.CITIES c
ON
    u.City = c.city
    AND u.State = c.state_name;
```

Utilizing the above view, we could use “Latitude” and “Longitude” information to populate the map visualization. Following is the output for our console-games data:



Section 1.2: Linear Regression on Price

Let's create a custom view to populate the required data from our database. Our goal is to run a linear regression model for price and using rating, user count, console count as independent variables.

We create following VIEW:

```
CREATE VIEW DB1SMOKE.VW_GameStats AS
WITH game_stats AS
(
SELECT
    g.GAMENAME ,
    g.PRICE ,
    g.RATING ,
    COUNT(s.USERID) AS UserCounts,
    COUNT(gc.CONSOLEID) AS ConsoleCounts
FROM
    DB1SMOKE.GAMES g
INNER JOIN DB1SMOKE.SUBSCRIPTIONS s
ON
    g.GAMEID = s.GAMEID
INNER JOIN DB1SMOKE.GAME_CONSOLES gc
ON
    gc.GAMEID = g.GAMEID
GROUP BY
    g.GAMENAME ,
    g.PRICE ,
    g.RATING
)
SELECT
    GameName,
    ROUND(AVG(Price), 3) AS Price,
    ROUND(AVG(Rating), 3) AS Rating,
    ROUND(AVG(UserCounts), 3) AS UserCounts,
    ROUND(AVG(ConsoleCounts), 3) AS ConsoleCounts
FROM
    game_stats
GROUP BY
    GameName
```

We export this data into a CSV file so that we can read it in a different statistical analysis tool. We initialized a linear regression model with following parameters:

- Dependent Variable: Price
- Independent Variables: Rating, Console Counts, Game Counts

Following is the output of the linear model:


```

Call:
lm(formula = Price ~ Rating + UserCounts + ConsoleCounts, data = data)

Residuals:
    Min       1Q   Median       3Q      Max
-27.3403  -4.4883  -0.2253   4.2104  27.9278

Coefficients: (1 not defined because of singularities)
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  48.064067   0.416564 115.382  <2e-16 ***
Rating       -1.370347   0.107295  -12.772  <2e-16 ***
UserCounts    0.001122   0.002422   0.463    0.643
ConsoleCounts NA         NA         NA         NA
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.522 on 14251 degrees of freedom
Multiple R-squared:  0.01132, Adjusted R-squared:  0.01118
F-statistic: 81.6 on 2 and 14251 DF, p-value: < 2.2e-16

```

We can observe that only Rating attribute shows the most significant factor in determining the price of the game. Number of Users playing, and Number of Console Counts has insignificant affect. This is as per the data in our database and might not represent the real-world scenarios.

Section 1.3: Initializing Machine Learning Models

Overview

In this experiment, our goal is to continue analyzing **Price** attribute. In one of the queries we created a **Cost Category** attribute that labels the Price attribute whether it is Very High, High, Average or Low. We will use this attribute to initialize two machine learning algorithms and compare the results.

Dataset Creation

We start by fetching the relevant data from our Database, to do this, we created a view that will be used to export this data into a CSV file:

```

CREATE VIEW [dbo].[VW_Game_Stats] AS
WITH stat AS
(
SELECT
    g.name AS GameName,
    g.price AS Price,
    g.rating AS Rating,

```

```

        count(s.userid) AS UserCounts,
        count(gc.consoleid) AS ConsoleCounts,
        CASE WHEN g.price >= 60 THEN 'Very High'
        WHEN g.price >=45 AND g.price < 60 THEN 'High'
        WHEN g.price >=30 AND g.price < 45 THEN 'Average'
        ELSE 'Low'
        END as CostCategory
FROM
    games g
INNER JOIN subscriptions s
ON
    g.gameid = s.gameid
INNER JOIN game_consoles gc
ON
    gc.gameid = g. gameid
GROUP BY
    g.name,
    g.price,
    g.rating
)
SELECT
    GameName,
    ROUND(AVG(Price), 3) AS Price,
    ROUND(AVG(Rating),3) AS Rating,
    ROUND(AVG(UserCounts),3) AS UserCounts,
    ROUND(AVG(ConsoleCounts),3) AS ConsoleCounts,
    CostCategory
FROM
    stat
GROUP BY GameName, CostCategory







```

Creating a ML Experiment

After creating the CSV, we import this dataset into **Azure ML Studio** (<https://studio.azureml.net/>). The dataset looks as shown in the Image below:


rows
14564

columns
6

| | GameName | Price | Rating | UserCounts | ConsoleCounts | CostCategory |
|---------|---|---|---|---|---|---|
| view as |  |  |  |  |  |  |
| | .38 Ambush Alley | 40 | 5 | 19 | 19 | Average |
| | 007 Legends | 39 | 3.72 | 21 | 21 | Average |
| | 007: A View to a Kill | 43 | 3.12 | 25 | 25 | Average |
| | 007: Licence to Kill | 44 | 3.12 | 100 | 100 | Average |
| | 007: Nightfire | 37 | 3.89 | 38 | 38 | Average |
| | 007: The Duel | 40 | 4.2 | 17 | 17 | Average |
| | 1 on 1 Government | 41 | 3.34 | 22 | 22 | Average |

The following charts shows the data distribution of the independent variables:


Histogram



Statistics

| | |
|--------------------|---------------|
| Mean | 42.921 |
| Median | 43 |
| Min | 15 |
| Max | 71 |
| Standard Deviation | 6.6305 |
| Unique Values | 87 |
| Missing Values | 0 |
| Feature Type | Numeric Featu |

Histogram



Statistics

| | |
|--------------------|-----------------|
| Mean | 3.7928 |
| Median | 3.82 |
| Min | 0.72 |
| Max | 5 |
| Standard Deviation | 0.5106 |
| Unique Values | 328 |
| Missing Values | 0 |
| Feature Type | Numeric Feature |

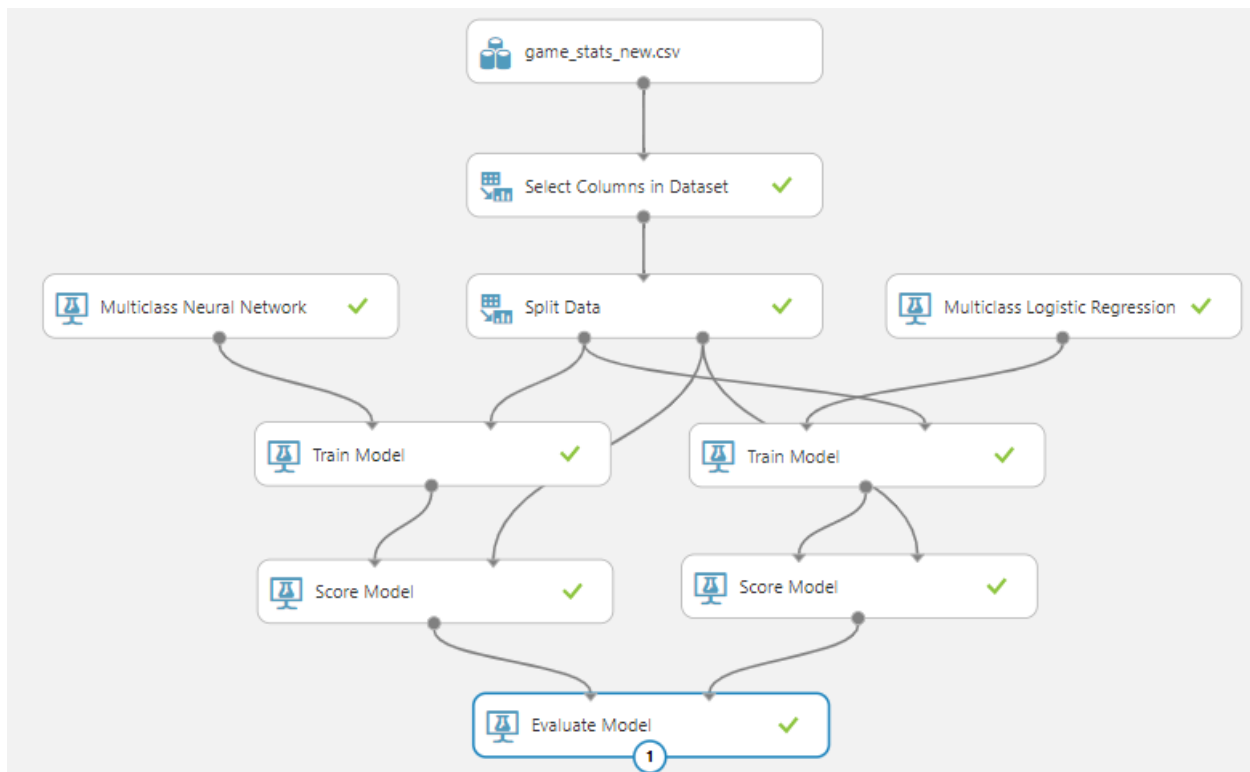
Initializing Model

We perform following steps in initializing model:

1. Import the dataset
2. Select the attributes to be used
 - a. In this case, the dependent variable is a Categorical Variable: **Cost Category**
 - b. The independent variables are **Price, Rating, Console Counts, Game Counts**
3. Split the train and test data by 70-30 split respectively.
4. Initialize Classification Algorithms
 - a. Multi-Class Neural Networks
 - b. Multi-Class Logistic Regression
5. Score the model using the test data and the Algorithm's outputs.
6. Evaluate the results for both for comparison.

We choose multi-Class algorithms because there are four classes viz: Very High, High, Average and Low. We choose classification algorithms because our dependent variable is a categorical variable.

The following figure the experiment created in Azure ML Studio:



Model Evaluation Results

DBMS - Experiment 5/1/2022 > Evaluate Model > Evaluation results

Metrics

| | |
|--------------------------|----------|
| Overall accuracy | 0.984894 |
| Average accuracy | 0.992447 |
| Micro-averaged precision | 0.984894 |
| Macro-averaged precision | 0.991287 |
| Micro-averaged recall | 0.984894 |
| Macro-averaged recall | 0.777618 |

Metrics

| | |
|--------------------------|----------|
| Overall accuracy | 0.972076 |
| Average accuracy | 0.986038 |
| Micro-averaged precision | 0.972076 |
| Macro-averaged precision | NaN |
| Micro-averaged recall | 0.972076 |
| Macro-averaged recall | 0.5 |


Left: Neural Networks | Right: Logistic Regression

Both the algorithms performed well with Neural Networks showing a slightly better accuracy.


Let's evaluate the confusion matrix, which is a very popular measure used while solving classification problems. It can be applied to binary classification as well as for multiclass classification problems.

The following image the confusion matrix comparison between the two algorithms:

Confusion Matrix



Confusion Matrix



Left: Neural Networks | Right: Logistic Regression

We observe that the Logistic Regression didn't perform well as the number of false-positives and false-negatives are quite high. The Neural Networks did better job in this regard.